
Stats202 Homework 3

By Fang Lin (Stanford ID # 06166564)

July 27, 2016

Problem 1 (p.260, ex3)

- (a)

Answer: Will steadily decreases. Since s increases from 0, training error for $s=0$ is the maximum and it keep decreasing to the Ordinary Least Square RSS.

- (b)

Answer: Decrease initially, and then eventually start increasing in a U shape. When $s=0$, all $\beta=0$, so the model is under-fitting. As increasing s , β becomes more fitting on the the test data, while after one point it will become over fitting.

- (c)

Answer: Steadily increase. Variance will keep increasing as long as the s increasing from 0.

- (d)

Answer: Steadily decrease. Bias will keep decreasing as long as the β inceasing from 0. When $s=0$, the model predict a constant so the bias is highest.

- (e)

Answer: Remain constant. Irreducible error is model independent.

Problem 2 (p.260, ex4)

- (a)

Answer: Steadily increase. As increasing numda from 0, beta decreases from least square estimate values to 0.

- (b)

Answer: Decrease initially, and then eventually start increasing in a U shape. When numda=0, all beta have their least square estimate values. So, in this case, the model can fit the training data best, but overfit cause high RSS. As numbda increasing, beta start decreasing to 0, so overfitting reduced, then as beta approach to 0, the model becomes too simple while test RSS will increase.

- (c)

Answer: Steadily decrease. Variance will keep decreasing as long as the numbda increasing from 0.

- (d)

Answer: Steadily increase. Bias will keep increasing as long as the numbda inceasing from 0. When numbda=0, the model has the least bias.

- (e)

Answer: Remain constant. Irreducible error is model independent.

Problem 3 (p.262, ex8)

- (a)

Answer:

```
> set.seed(1)
> X = rnorm(100)
> eps = rnorm(100)
```

- (b)

Answer:


We select $\beta_0=1$, $\beta_1=2$, $\beta_2=-1$, $\beta_3=0.5$.

```
> beta0=1
> beta1=2
> beta2=-1
> beta3=0.5
> Y = beta0 + beta1 * X + beta2 * X^2 + beta3 * X^3 + eps
```

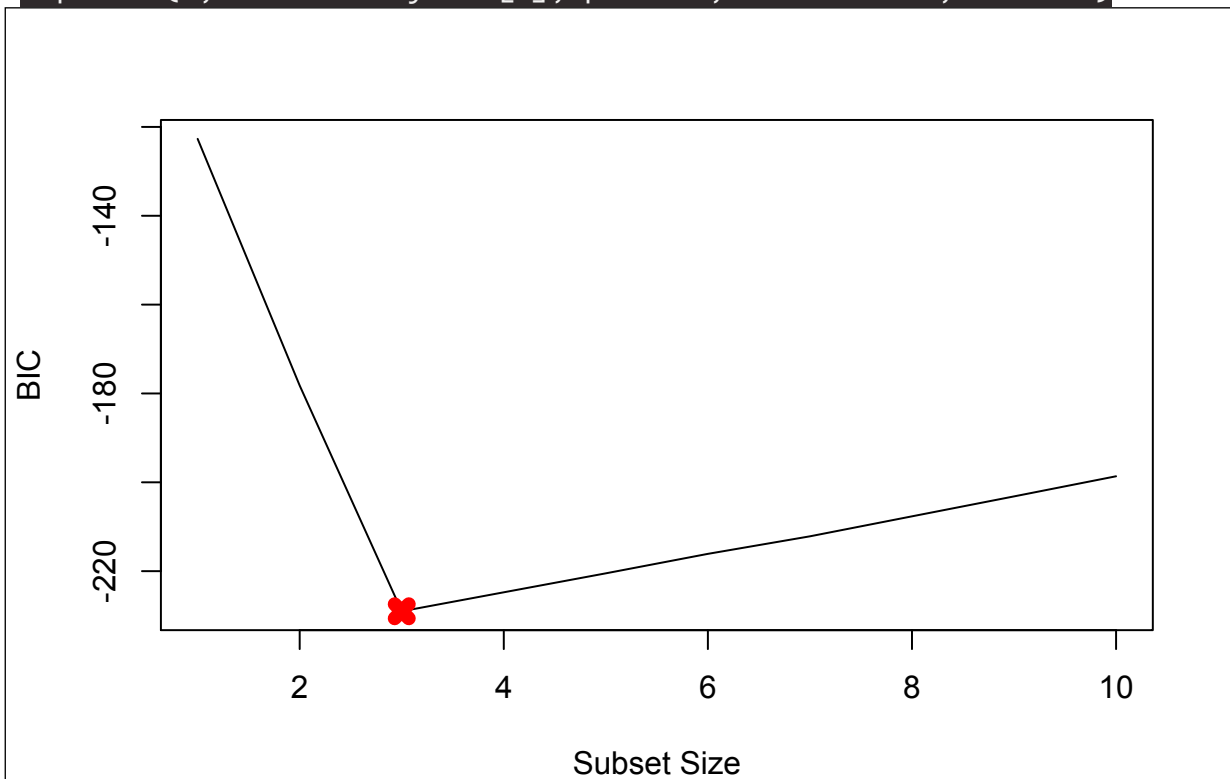
- (c)

Answer:

```
> install.packages("leaps")
> library(leaps)
> data.full = data.frame(y = Y, x = X)
> mod.full = regsubsets(y ~ poly(x, 10, raw = T), data = data.full,
nvmax = 10)
> mod.summary = summary(mod.full)
> which.min(mod.summary$cp)
[1] 3
> which.min(mod.summary$bic)
[1] 3
> which.max(mod.summary$adjr2)
[1] 3
> plot(mod.summary$cp, xlab = "Subset Size", ylab = "Cp", pch = 20,
type = "l")
> points(3, mod.summary$cp[3], pch = 4, col = "red", lwd = 7)
```




```
> plot(mod.summary$bic, xlab = "Subset Size", ylab = "BIC", pch = 20,
type = "l")
> points(3, mod.summary$bic[3], pch = 4, col = "red", lwd = 7)
```



```
> plot(mod.summary$adjr2, xlab = "Subset Size", ylab = "Adjusted R2",
pch = 20,
```

```
+ type = "l")
> points(3, mod.summary$adjr2[3], pch = 4, col = "red", lwd = 7)
```



```
> plot(mod.summary$adjr2, xlab = "Subset Size", ylab = "Adjusted R2",
+ type = "l")
> points(3, mod.summary$adjr2[3], pch = 4, col = "red", lwd = 7)
> coefficients(mod.full, id = 3)
(Intercept) poly(x, 10, raw = T)1 poly(x, 10, raw = T)2
1.07219472 2.44514720 -1.15676236
poly(x, 10, raw = T)5
0.09022577
```

Answer: With Cp, BIC and Adjusted R2 criteria, 3, 3, 3 variable models are picked. All statistics pick X^5 over X^3 . The remaining coefficients are quite close to betas.

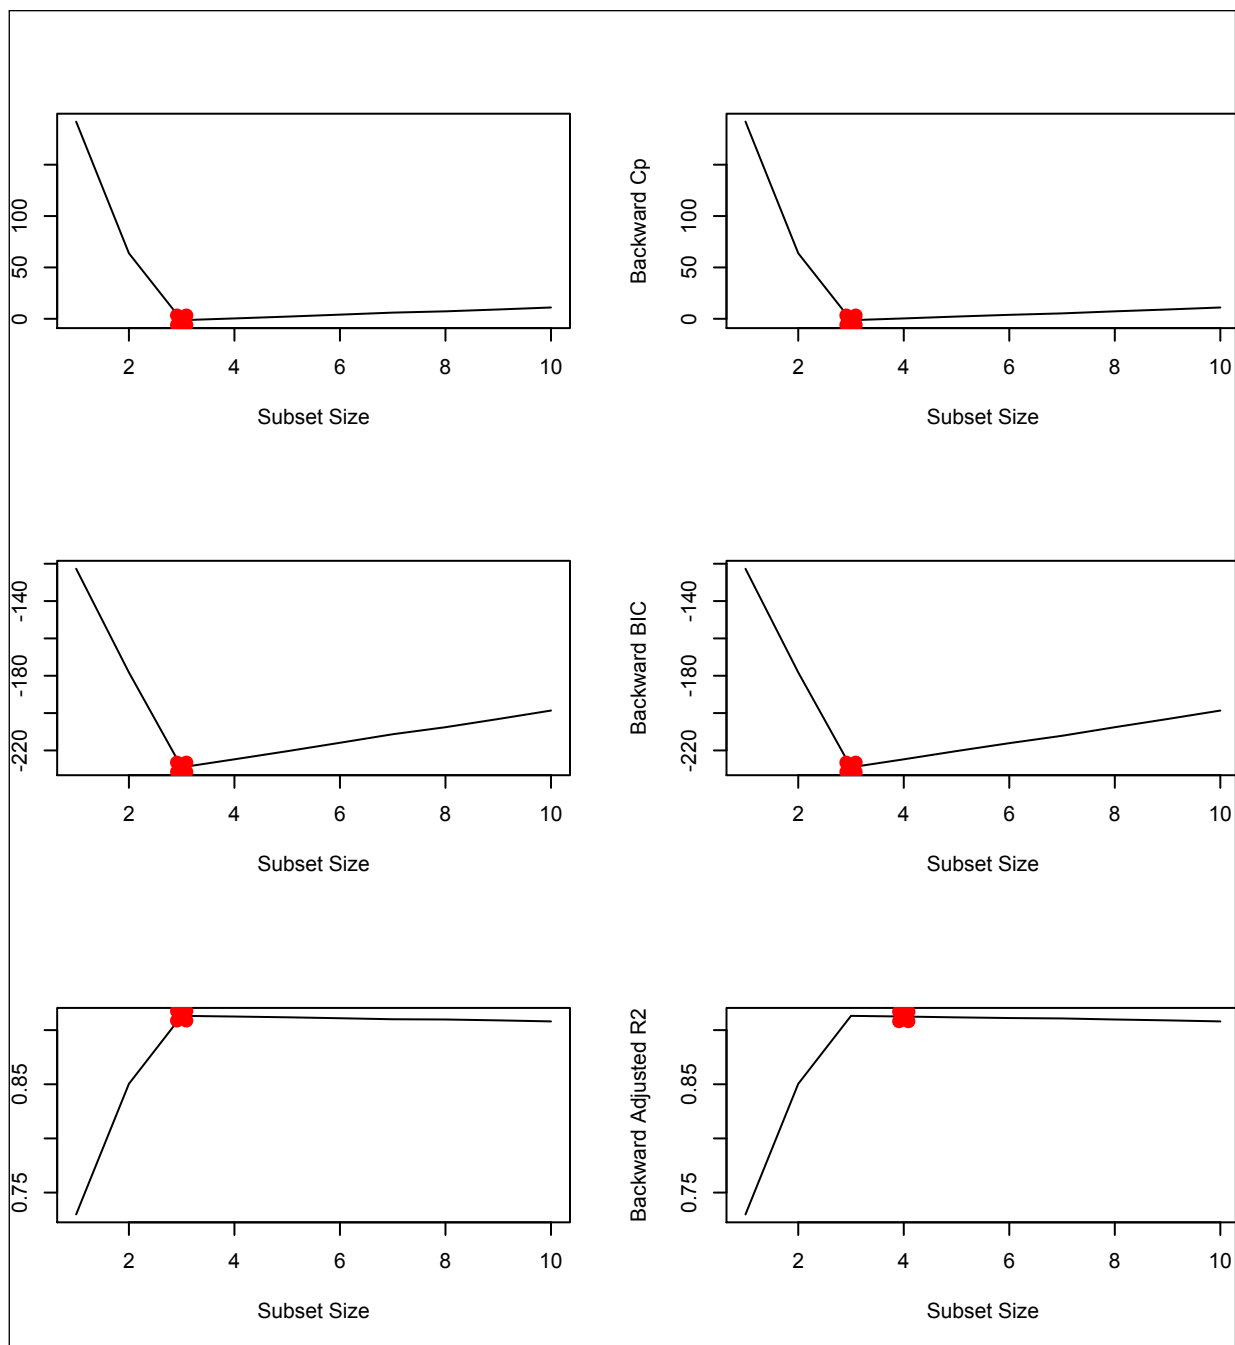
- (d)

```
> mod.fwd = regsubsets(y ~ poly(x, 10, raw = T), data = data.full,
+ nvmax = 10,
+ method = "forward")
> mod.bwd = regsubsets(y ~ poly(x, 10, raw = T), data = data.full,
+ nvmax = 10,
+ method = "backward")
> fwd.summary = summary(mod.fwd)
> bwd.summary = summary(mod.bwd)
> which.min(fwd.summary$cp)
```

```

[1] 3
> which.min(bwd.summary$cp)
[1] 3
>
> which.min(fwd.summary$bic)
[1] 3
> which.min(bwd.summary$bic)
[1] 3
> which.max(fwd.summary$adjr2)
[1] 3
> which.max(bwd.summary$adjr2)
[1] 3
> par(mfrow = c(3, 2))
> plot(fwd.summary$cp, xlab = "Subset Size", ylab = "Forward Cp", pch
= 20, type = "l")
> points(3, fwd.summary$cp[3], pch = 4, col = "red", lwd = 7)
> plot(bwd.summary$cp, xlab = "Subset Size", ylab = "Backward Cp",
pch = 20, type = "l")
> points(3, bwd.summary$cp[3], pch = 4, col = "red", lwd = 7)
> plot(fwd.summary$bic, xlab = "Subset Size", ylab = "Forward BIC",
pch = 20,
+   type = "l")
> points(3, fwd.summary$bic[3], pch = 4, col = "red", lwd = 7)
> plot(bwd.summary$bic, xlab = "Subset Size", ylab = "Backward BIC",
pch = 20,
+   type = "l")
> points(3, bwd.summary$bic[3], pch = 4, col = "red", lwd = 7)
> plot(fwd.summary$adjr2, xlab = "Subset Size", ylab = "Forward
Adjusted R2",
+   pch = 20, type = "l")
> points(3, fwd.summary$adjr2[3], pch = 4, col = "red", lwd = 7)
> plot(bwd.summary$adjr2, xlab = "Subset Size", ylab = "Backward
Adjusted R2",
+   pch = 20, type = "l")
> points(4, bwd.summary$adjr2[4], pch = 4, col = "red", lwd = 7)

```



```
> coefficients(mod.fwd, id = 3)
(Intercept) poly(x, 10, raw = T)1 poly(x, 10, raw = T)2
1.07219472 2.44514720 -1.15676236
poly(x, 10, raw = T)5
0.09022577
> coefficients(mod.bwd, id = 3)
(Intercept) poly(x, 10, raw = T)1 poly(x, 10, raw = T)2
1.07219472 2.44514720 -1.15676236
poly(x, 10, raw = T)5
0.09022577
```



```

0.09022577
> coefficients(mod.fwd, id = 4)
      (Intercept) poly(x, 10, raw = T)1 poly(x, 10, raw = T)2
      1.11309290      2.46288862      -1.28112227
poly(x, 10, raw = T)4 poly(x, 10, raw = T)5
      0.03074107      0.08769746

```

Answer: Forward stepwise picks X^5 over X^3 . Backward stepwise with 3 variable picks X^5 while backward stepwise with 4 variables picks X^4 and X^5 . All other coefficients are close to betas.


- (e)

```

> library(glmnet)
Loading required package: Matrix
Loading required package: foreach
foreach: simple, scalable parallel programming from Revolution
Analytics
Use Revolution R for scalability, fault tolerance and more.
http://www.revolutionanalytics.com
Loaded glmnet 2.0-5

> xmat = model.matrix(y ~ poly(x, 10, raw = T), data = data.full)[,
-1]
> mod.lasso = cv.glmnet(xmat, Y, alpha = 1)
> best.lambda = mod.lasso$lambda.min
> best.lambda
[1] 0.03779912
> plot(mod.lasso)

```



```
> best.model = glmnet(xmat, Y, alpha = 1)
> predict(best.model, s = best.lambda, type = "coefficients")
11 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept) 1.04103683
poly(x, 10, raw = T)1 2.28787832
poly(x, 10, raw = T)2 -1.10681293
poly(x, 10, raw = T)3 0.13650159
poly(x, 10, raw = T)4 .
poly(x, 10, raw = T)5 0.06464279
poly(x, 10, raw = T)6 .
poly(x, 10, raw = T)7 .
poly(x, 10, raw = T)8 .
poly(x, 10, raw = T)9 .
poly(x, 10, raw = T)10 .
```

Answer: Lasso picks X^3 also little bit X^5 .

- (e)

```
> beta7 = 7
> Y = beta0 + beta7 * X^7 + eps
```

```

> # Predict using regsubsets
> data.full = data.frame(y = Y, x = X)
> mod.full = regsubsets(y ~ poly(x, 10, raw = T), data = data.full,
nvmax = 10)
> mod.summary = summary(mod.full)
>
> # Find the model size for best cp, BIC and adjr2
> which.min(mod.summary$cp)
[1] 2
> which.min(mod.summary$bic)
[1] 1
> which.max(mod.summary$adjr2)
[1] 4
> coefficients(mod.full, id = 1)
      (Intercept) poly(x, 10, raw = T)7
      0.9589402      7.0007705
> coefficients(mod.full, id = 2)
      (Intercept) poly(x, 10, raw = T)2 poly(x, 10, raw = T)7
      1.0704904      -0.1417084      7.0015552
> coefficients(mod.full, id = 4)
      (Intercept) poly(x, 10, raw = T)1 poly(x, 10, raw = T)2
      1.0762524      0.2914016      -0.1617671
poly(x, 10, raw = T)3 poly(x, 10, raw = T)7
      -0.2526527      7.0091338
> xmat = model.matrix(y ~ poly(x, 10, raw = T), data = data.full)[,
-1]
> mod.lasso = cv.glmnet(xmat, Y, alpha = 1)
> best.lambda = mod.lasso$lambda.min
> best.lambda
[1] 13.57478
> best.model = glmnet(xmat, Y, alpha = 1)
> predict(best.model, s = best.lambda, type = "coefficients")
11 x 1 sparse Matrix of class "dgCMatrix"
      1
(Intercept)      1.904188
poly(x, 10, raw = T)1 .
poly(x, 10, raw = T)2 .
poly(x, 10, raw = T)3 .
poly(x, 10, raw = T)4 .
poly(x, 10, raw = T)5 .
poly(x, 10, raw = T)6 .
poly(x, 10, raw = T)7 6.776797
poly(x, 10, raw = T)8 .

```

```
poly(x, 10, raw = T)9 .  
poly(x, 10, raw = T)10 .
```

Answer: We see BIC and Lasso both pick the best 1-variable.

Problem 4 (p.263, ex9)

- (a)

```
> library(ISLR)
> set.seed(11)
> sum(is.na(College))
[1] 0
> train.size = dim(College)[1] / 2
> train = sample(1:dim(College)[1], train.size)
> test = -train
> College.train = College[train, ]
> College.test = College[test, ]
```

- (b)

```
> lm.fit = lm(Apps~., data=College.train)
> lm.pred = predict(lm.fit, College.test)
> mean((College.test[, "Apps"] - lm.pred)^2)
[1] 1538442
```

- (c)

```
> library(glmnet)
> train.mat = model.matrix(Apps~., data=College.train)
> test.mat = model.matrix(Apps~., data=College.test)
> grid = 10 ^ seq(4, -2, length=100)
> mod.ridge = cv.glmnet(train.mat, College.train[, "Apps"], alpha=0,
lambda=grid, thresh=1e-12)
> lambda.best = mod.ridge$lambda.min
> lambda.best
[1] 18.73817
> ridge.pred = predict(mod.ridge, newx=test.mat, s=lambda.best)
> mean((College.test[, "Apps"] - ridge.pred)^2)
[1] 1608859
```

Answer: The RSS is slightly higher than OLS, 1608859.

- (d)

```
> mod.lasso = cv.glmnet(train.mat, College.train[, "Apps"], alpha=1,
lambda=grid, thresh=1e-12)
> lambda.best = mod.lasso$lambda.min
> lambda.best
[1] 21.54435
> lasso.pred = predict(mod.lasso, newx=test.mat, s=lambda.best)
> mean((College.test[, "Apps"] - lasso.pred)^2)
[1] 1635280
```

- (e)

```
> library(pls)
```

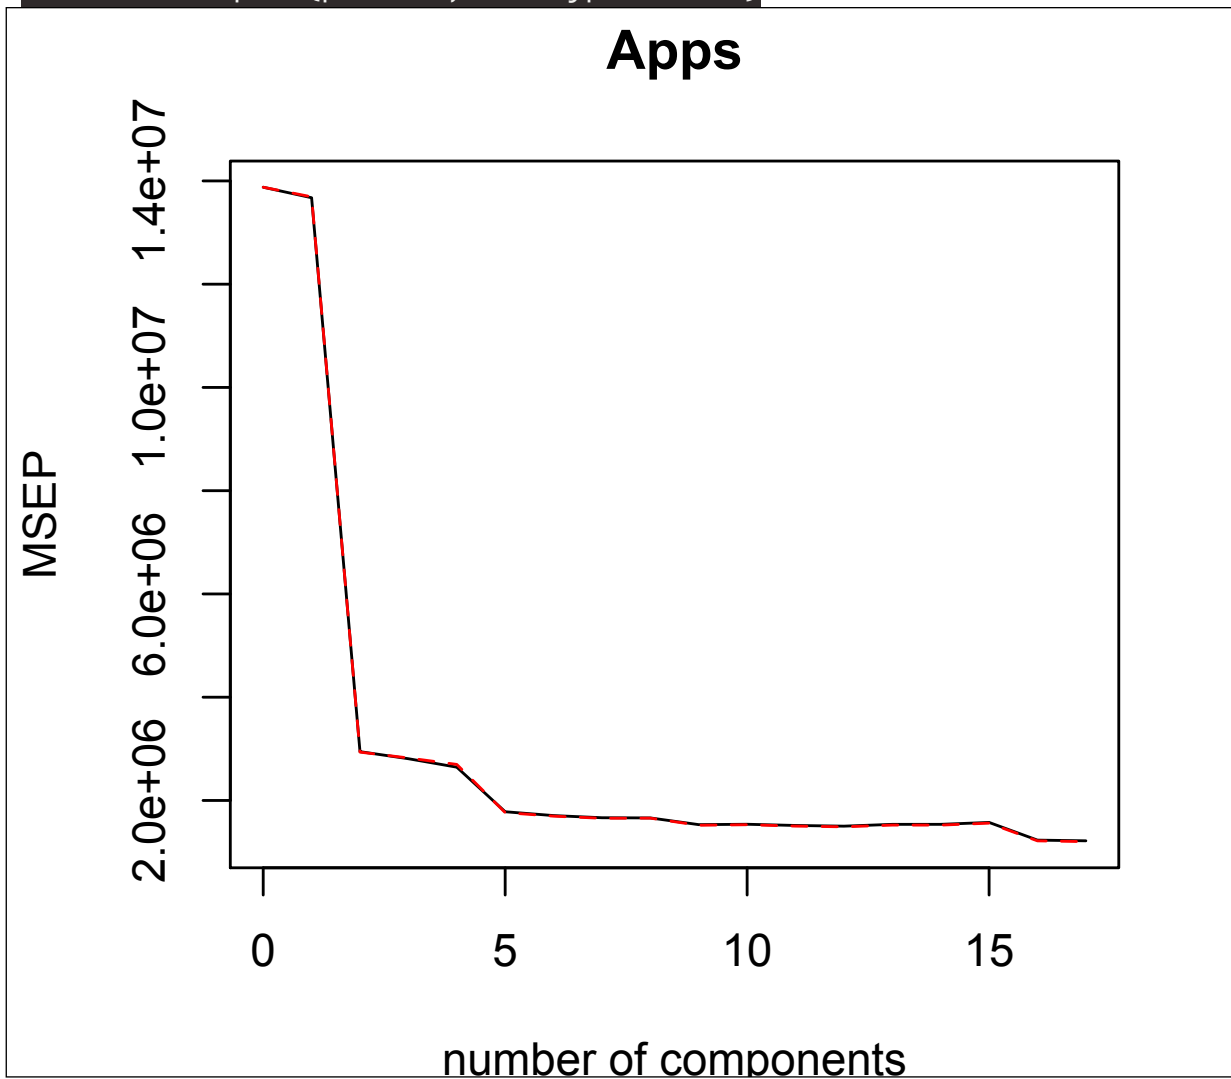
Attaching package: 'pls'

The following object is masked from 'package:stats':

loadings

```
> pcr.fit = pcr(Apps~., data=College.train, scale=T, validation="CV")
```

```
> validationplot(pcr.fit, val.type="MSEP")
```



```
> pcr.pred = predict(pcr.fit, College.test, ncomp=10)
```


```
> mean((College.test[, "Apps"] - data.frame(pcr.pred))^2)
```

```
[1] 3014496
```

- (f)

```
> pls.fit = plsr(Apps~., data=College.train, scale=T,  
validation="CV")
```

```
> validationplot(pls.fit, val.type="MSEP")
```

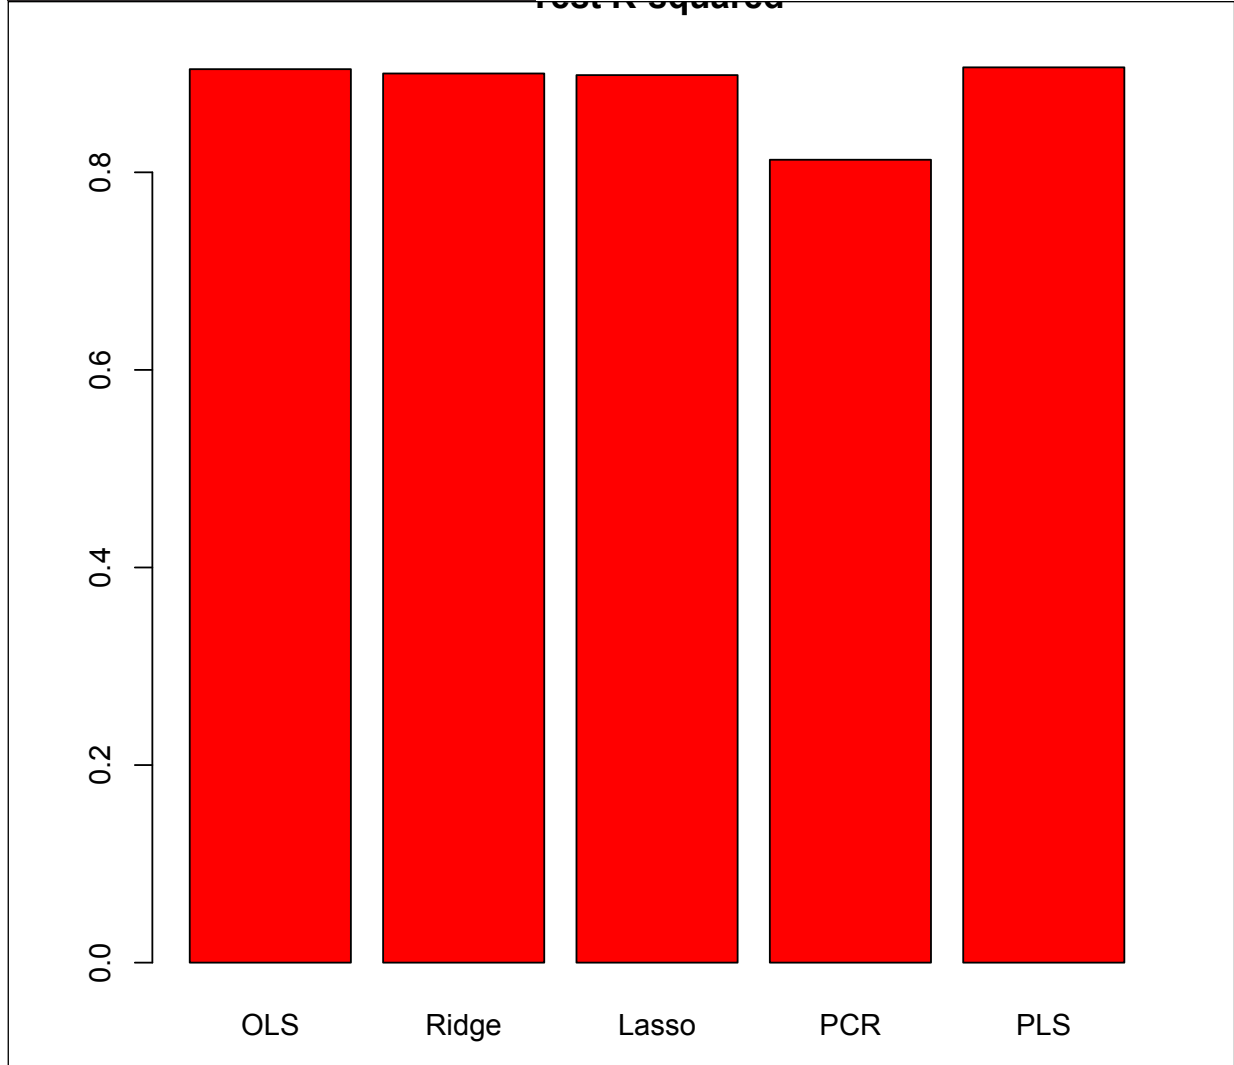


```
> pls.pred = predict(pls.fit, College.test, ncomp=10)
> mean((College.test[, "Apps"] - data.frame(pls.pred))^2)
[1] 1508987
```

- (g)

```
> test.avg = mean(College.test[, "Apps"])
> lm.test.r2 = 1 - mean((College.test[, "Apps"] - lm.pred)^2) /
mean((College.test[, "Apps"] - test.avg)^2)
> ridge.test.r2 = 1 - mean((College.test[, "Apps"] - ridge.pred)^2) /
mean((College.test[, "Apps"] - test.avg)^2)
> lasso.test.r2 = 1 - mean((College.test[, "Apps"] - lasso.pred)^2) /
mean((College.test[, "Apps"] - test.avg)^2)
> pcr.test.r2 = 1 - mean((College.test[, "Apps"] -
data.frame(pcr.pred))^2) / mean((College.test[, "Apps"] - test.avg)^2)
> pls.test.r2 = 1 - mean((College.test[, "Apps"] -
data.frame(pls.pred))^2) / mean((College.test[, "Apps"] - test.avg)^2)
```

```
> barplot(c(lm.test.r2, ridge.test.r2, lasso.test.r2, pcr.test.r2,
pls.test.r2), col="red", names.arg=c("OLS", "Ridge", "Lasso", "PCR",
"PLS"), main="Test R-squared")
```




Answer: The plot shows test R^2 for different models other than PCR are near 0.9. PCR has a smaller test square root. All models except PCR predict college applications with high accuracy.

Problem 5 (p.332, ex4)

(a)


Answer:



(b)

```
> par(xpd = NA)
> plot(NA, NA, type = "n", xlim = c(-2, 2), ylim = c(-3, 3), xlab =
"X1", ylab = "X2")
> # X2 < 1
> lines(x = c(-2, 2), y = c(1, 1))
> # X1 < 1 with X2 < 1
> lines(x = c(1, 1), y = c(-3, 1))
> text(x = (-2 + 1)/2, y = -1, labels = c(-1.8))
> text(x = 1.5, y = -1, labels = c(0.63))
> # X2 < 2 with X2 >= 1
> lines(x = c(-2, 2), y = c(2, 2))
> text(x = 0, y = 2.5, labels = c(2.49))
> # X1 < 0 with X2 < 2 and X2 >= 1
> lines(x = c(0, 0), y = c(1, 2))
> text(x = -1, y = 1.5, labels = c(-1.06))
> text(x = 1, y = 1.5, labels = c(0.21))
```

Answer:



Problem 6 (p.332, ex5)

(a)

```
> p = c(0.1, 0.15, 0.2, 0.2, 0.55, 0.6, 0.6, 0.65, 0.7, 0.75)
> sum(p >= 0.5) > sum(p < 0.5)
[1] TRUE
> mean(p)
[1] 0.45
```

Answer:

Using 50% threshold, the number of red predictions is larger than the number of green predictions. Average approach will choose green, since the average of the probabilities is less than the 50 threshold.

Problem 7 (p.333, ex8)

(a)

Answer:

```
> library(ISLR)
> attach(Carseats)
> set.seed(1)
> 
> train = sample(dim(Carseats)[1], dim(Carseats)[1]/2)
> Carseats.train = Carseats[train, ]
> Carseats.test = Carseats[-train, ]
```

(b)

```
> library(tree)
> tree.carseats = tree(Sales ~ ., data = Carseats.train)
> summary(tree.carseats)
```

Regression tree:

```
tree(formula = Sales ~ ., data = Carseats.train)
```

Variables actually used in tree construction:

```
[1] "ShelveLoc" "Price" "Age" "Advertising" "Income"
[6] "CompPrice"
```

Number of terminal nodes: 18

Residual mean deviance: 2.36 = 429.5 / 182

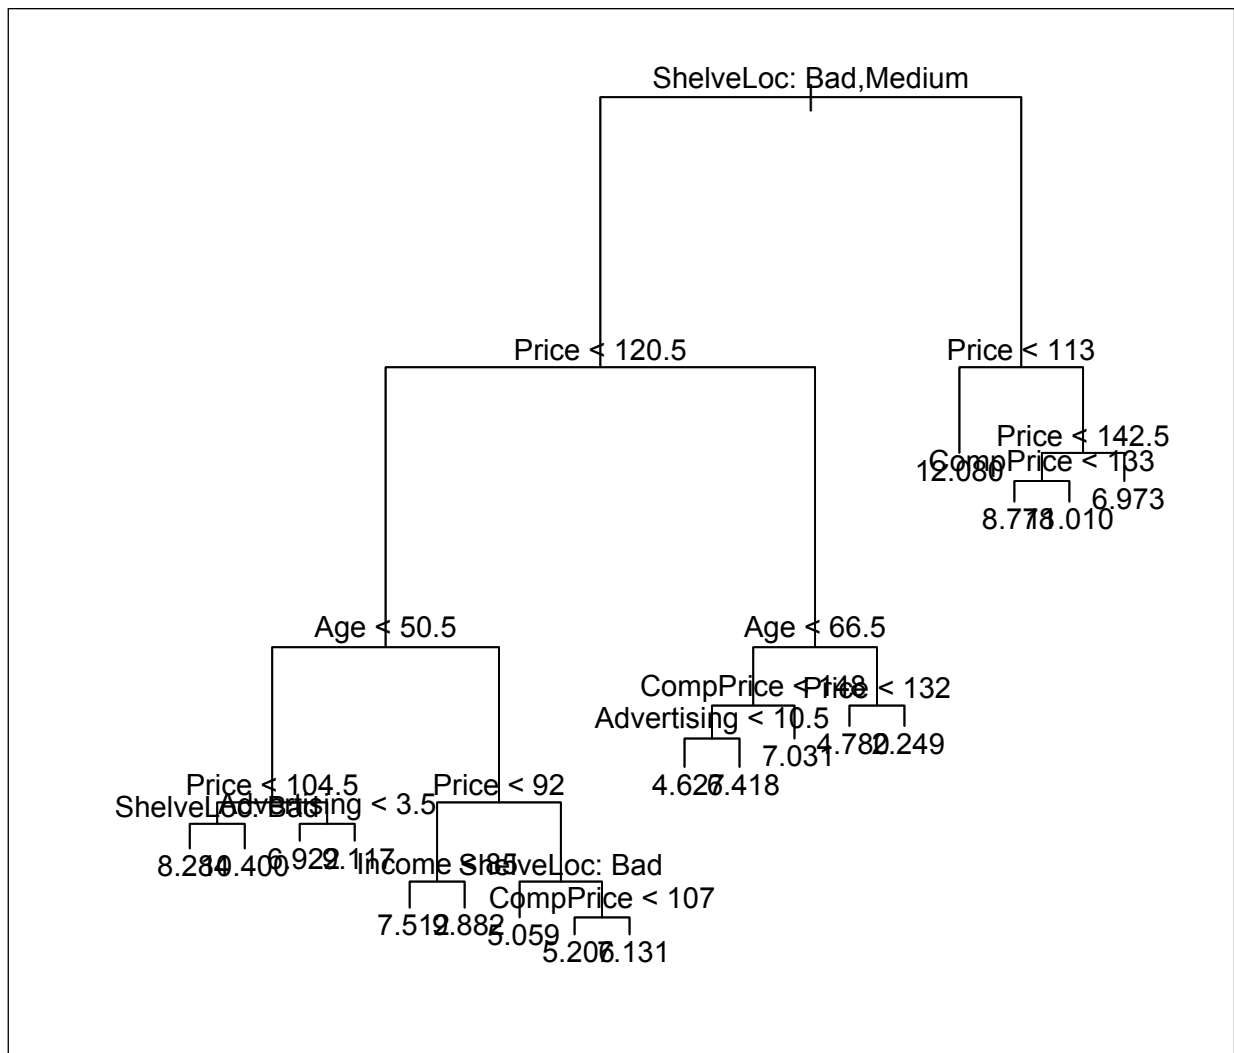
Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-4.2570	-1.0360	0.1024	0.0000	0.9301	3.9130

```
> plot(tree.carseats)
```

```
> text(tree.carseats, pretty = 0)
```

Answer:



```

> pred.carseats = predict(tree.carseats, Carseats.test)
> mean((Carseats.test$Sales - pred.carseats)^2)
[1] 4.148897

```

(c)

```

> cv.carseats = cv.tree(tree.carseats, FUN = prune.tree)
> par(mfrow = c(1, 2))
> # Best size = 9
> pruned.carseats = prune.tree(tree.carseats, best = 9)
> par(mfrow = c(1, 1))
> plot(pruned.carseats)
> text(pruned.carseats, pretty = 0):

```

Answer:

```

> pred.pruned = predict(pruned.carseats, Carseats.test)
> mean((Carseats.test$Sales - pred.pruned)^2)
[1] 4.993124

```

Answer: No, increased instead.

(d)

```
> library(randomForest)
randomForest 4.6-12
Type rfNews() to see new features/changes/bug fixes.
> bag.carseats = randomForest(Sales ~ ., data = Carseats.train, mtry
= 10, ntree = 500,
+   importance = T)
> bag.pred = predict(bag.carseats, Carseats.test)
> mean((Carseats.test$Sales - bag.pred)^2)
[1] 2.604369
> importance(bag.carseats)
```

	%IncMSE	IncNodePurity
CompPrice	14.4124562	133.731797
Income	6.5147532	74.346961
Advertising	15.7607104	117.822651
Population	0.6031237	60.227867
Price	57.8206926	514.802084
ShelveLoc	43.0486065	319.117972
Age	19.8789659	192.880596
Education	2.9319161	39.490093
Urban	-3.1300102	8.695529
US	7.6298722	15.723975

Answer:

Bagging improved the test MSE to 2.6. Price, ShelveLoc and Age are most important predictors of Sale.

(e)

```
> rf.carseats = randomForest(Sales ~ ., data = Carseats.train, mtry =
5, ntree = 500,
+   importance = T)
> rf.pred = predict(rf.carseats, Carseats.test)
mean((Carseats.test$Sales - rf.pred)^2)
> mean((Carseats.test$Sales - rf.pred)^2)
[1] 2.802383
> importance(rf.carseats)
```

	%IncMSE	IncNodePurity
CompPrice	12.0259791	124.81403
Income	5.5542673	106.15418
Advertising	12.0466048	136.15204
Population	0.3136897	81.68162
Price	45.9639857	457.15711

ShelveLoc	36.2789679	271.76488
Age	20.8537727	196.72182
Education	2.9005332	54.16980
Urban	-0.6888196	11.86848
US	6.9739759	23.64075.

Answer:

Random forest worsens the MSE on the test set to 2.80. Price, ShelveLoc and Age are three most important predictors of Sale.