
Stats202 Homework 4

By Fang Lin (Stanford ID # 06166564)

July 30, 2016

Problem 1 (p.334, ex10)

- (a)

Answer:

```
> library(ISLR)
> sum(is.na(Hitters$Salary))
[1] 59
> Hitters = Hitters[-which(is.na(Hitters$Salary)), ]
> sum(is.na(Hitters$Salary))
[1] 0
> Hitters$Salary = log(Hitters$Salary)
```

- (b)

Answer:

```
> train = 1:200
> Hitters.train = Hitters[train, ]
> Hitters.test = Hitters[-train, ]
```

- (c)

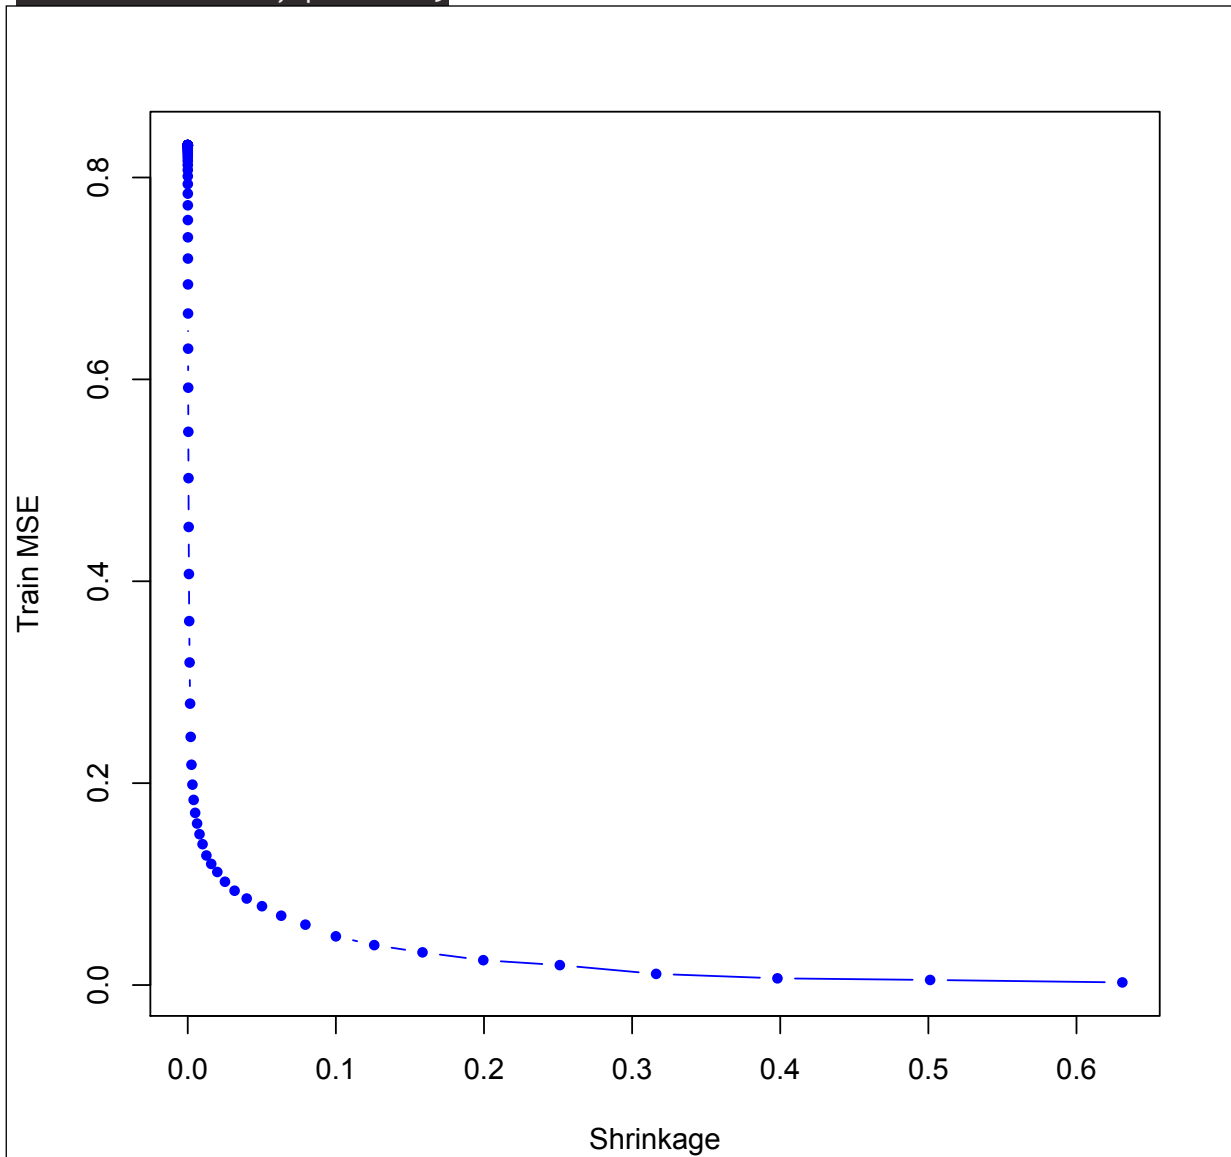
Answer: >

library(gbm)

```
Loading required package: survival
Loading required package: lattice
Loading required package: splines
Loading required package: parallel
Loaded gbm 2.1.1
> set.seed(103)
> pows = seq(-10, -0.2, by = 0.1)
> lambdas = 10^pows
> length.lambdas = length(lambdas)
> train.errors = rep(NA, length.lambdas)
> test.errors = rep(NA, length.lambdas)
> for (i in 1:length.lambdas) {
+   boost.hitters = gbm(Salary ~ ., data = Hitters.train, distribution =
"gaussian",
+   n.trees = 1000, shrinkage = lambdas[i])
+   train.pred = predict(boost.hitters, Hitters.train, n.trees = 1000)
+   test.pred = predict(boost.hitters, Hitters.test, n.trees = 1000)
+   train.errors[i] = mean((Hitters.train$Salary - train.pred)^2)
+   test.errors[i] = mean((Hitters.test$Salary - test.pred)^2)
+ }

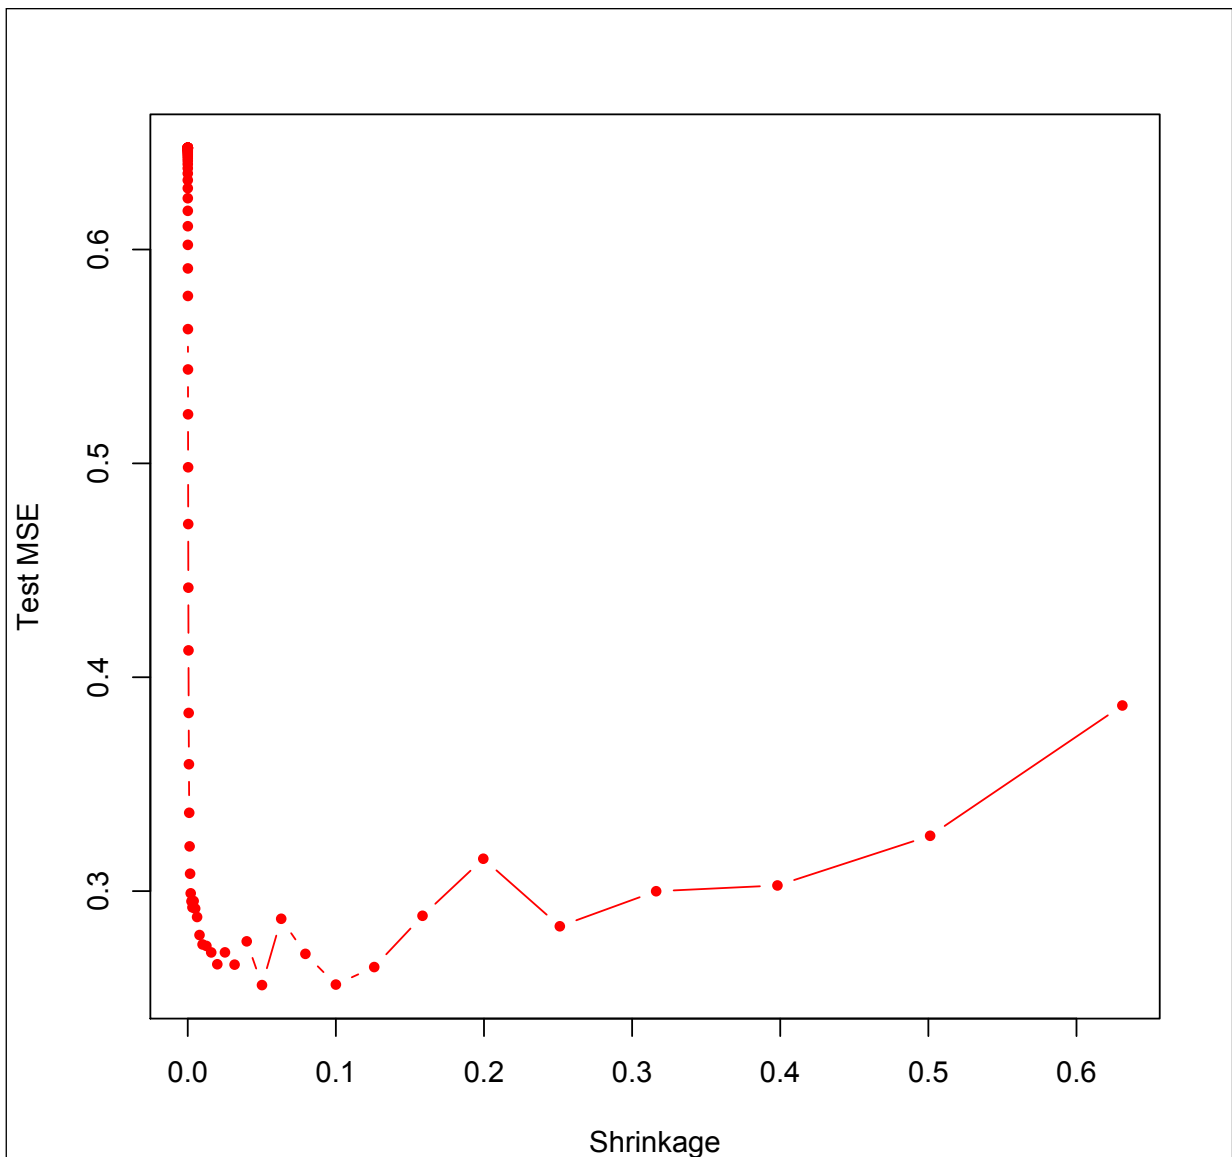
plot(lambdas, train.errors, type = "b", xlab = "Shrinkage", ylab = "Train
MSE",
     col = "blue", pch = 20)
```

```
> plot(lambdas, train.errors, type = "b", xlab = "Shrinkage", ylab = "Train  
MSE",  
+ col = "blue", pch = 20)
```



- (d)

Answer:



```
> min(test.errors)
```

```
[1] 0.2560507
```

```
> lambdas[which.min(test.errors)]
```

```
[1] 0.05011872
```

When numbd=0.05011872, the minium test error got 0.2561

- (e)

Answer:

```
> lm.fit = lm(Salary ~ ., data = Hitters.train)
```

```
> lm.pred = predict(lm.fit, Hitters.test)
```

```
> mean((Hitters.test$Salary - lm.pred)^2)
```

```
[1] 0.4917959
```

```
> library(glmnet)
```

```
Loading required package: Matrix
```

```
Loading required package: foreach
```

foreach: simple, scalable parallel programming from Revolution Analytics
Use Revolution R for scalability, fault tolerance and more.
<http://www.revolutionanalytics.com>
Loaded glmnet 2.0-5

```
> set.seed(134)
> x = model.matrix(Salary ~ ., data = Hitters.train)
> y = Hitters.train$Salary
> x.test = model.matrix(Salary ~ ., data = Hitters.test)
> lasso.fit = glmnet(x, y, alpha = 1)
> lasso.pred = predict(lasso.fit, s = 0.01, newx = x.test)
> mean((Hitters.test$Salary - lasso.pred)^2)
[1] 0.4700537
```

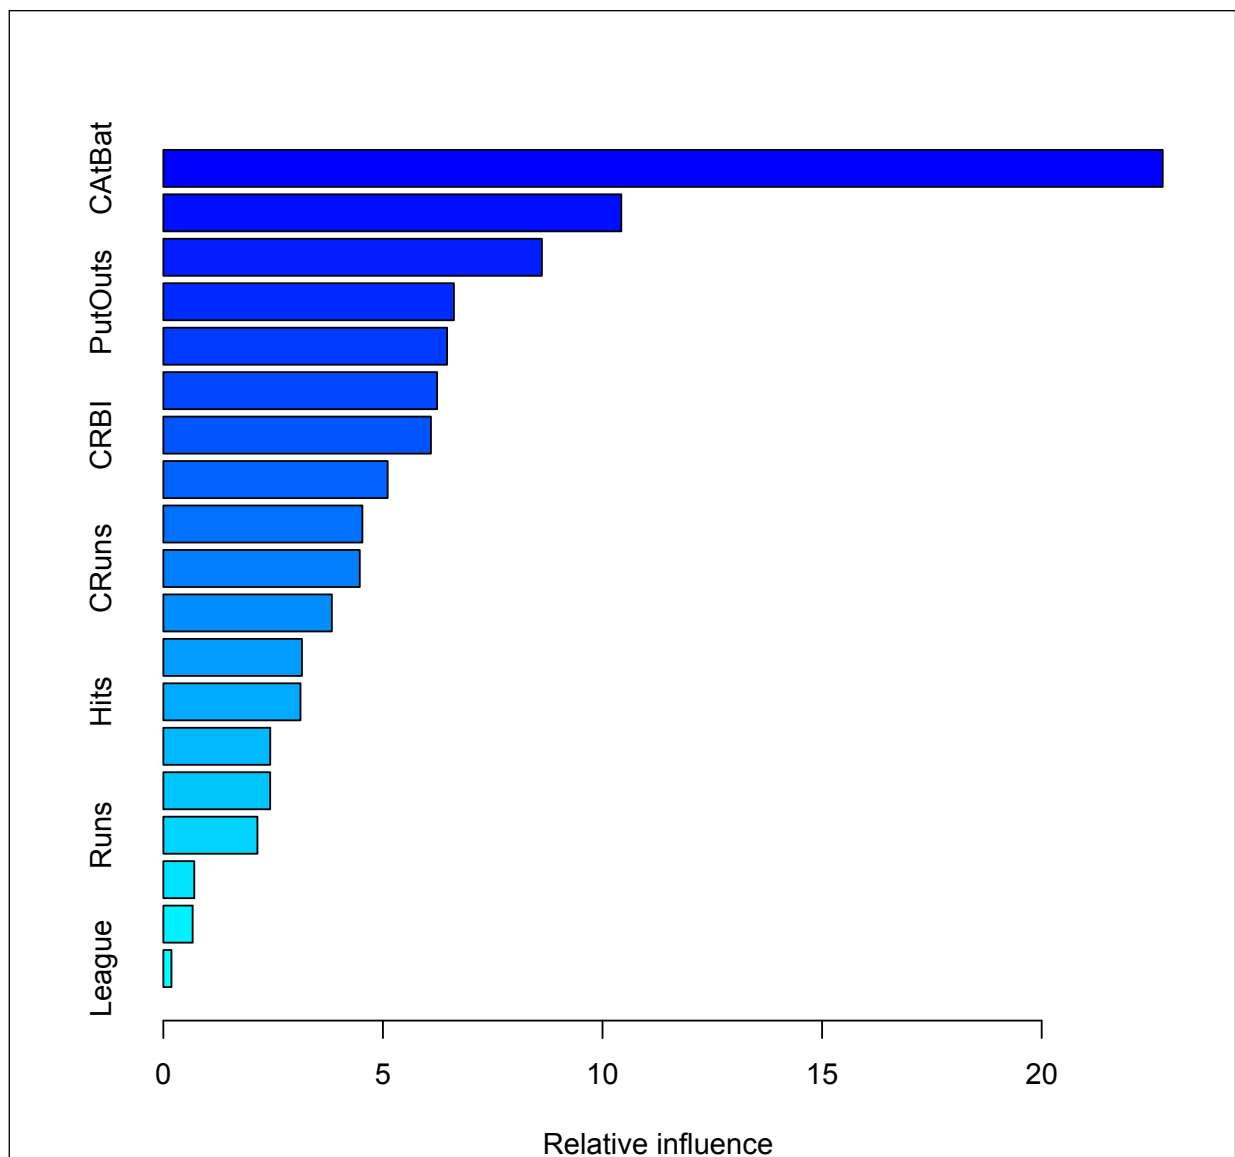
Both linear and regularization models have higher test MSE than boosting.

- (f)

Answer:

```
> boost.best = gbm(Salary ~ ., data = Hitters.train, distribution =
"gaussian",
+   n.trees = 1000, shrinkage = lambdas[which.min(test.errors)])
> summary(boost.best)
```

| | var | rel.inf |
|-----------|-----------|------------|
| CAtBat | CAtBat | 22.7562681 |
| CWalks | CWalks | 10.4279674 |
| CHits | CHits | 8.6198109 |
| PutOuts | PutOuts | 6.6159325 |
| Years | Years | 6.4611683 |
| Walks | Walks | 6.2331148 |
| CRBI | CRBI | 6.0926744 |
| CHmRun | CHmRun | 5.1076104 |
| RBI | RBI | 4.5321678 |
| CRuns | CRuns | 4.4728132 |
| Assists | Assists | 3.8366575 |
| HmRun | HmRun | 3.1554038 |
| Hits | Hits | 3.1229284 |
| AtBat | AtBat | 2.4338530 |
| Errors | Errors | 2.4324185 |
| Runs | Runs | 2.1425481 |
| Division | Division | 0.7041949 |
| NewLeague | NewLeague | 0.6675446 |
| League | League | 0.1849234 |



CRBI, CWalks and CATBat are the three most important variables.

- (g)

Answer:

```
> library(randomForest)
randomForest 4.6-12
Type rfNews() to see new features/changes/bug fixes.
> set.seed(21)
> rf.hitters = randomForest(Salary ~ ., data = Hitters.train, ntree =
500, mtry = 19)
> rf.pred = predict(rf.hitters, Hitters.test)
> mean((Hitters.test$Salary - rf.pred)^2)
[1] 0.231884.
```

Bagging produce 0.23 Test MSE, which is lower than the best test MSE of boosting.

Problem 2 (p.335, ex11)

- (a)

Answer:

```
> library(ISLR)
> train = 1:1000
> Caravan$Purchase = ifelse(Caravan$Purchase == "Yes", 1, 0)
> Caravan.train = Caravan[train, ]
> Caravan.test = Caravan[-train, ]
```

- (b)

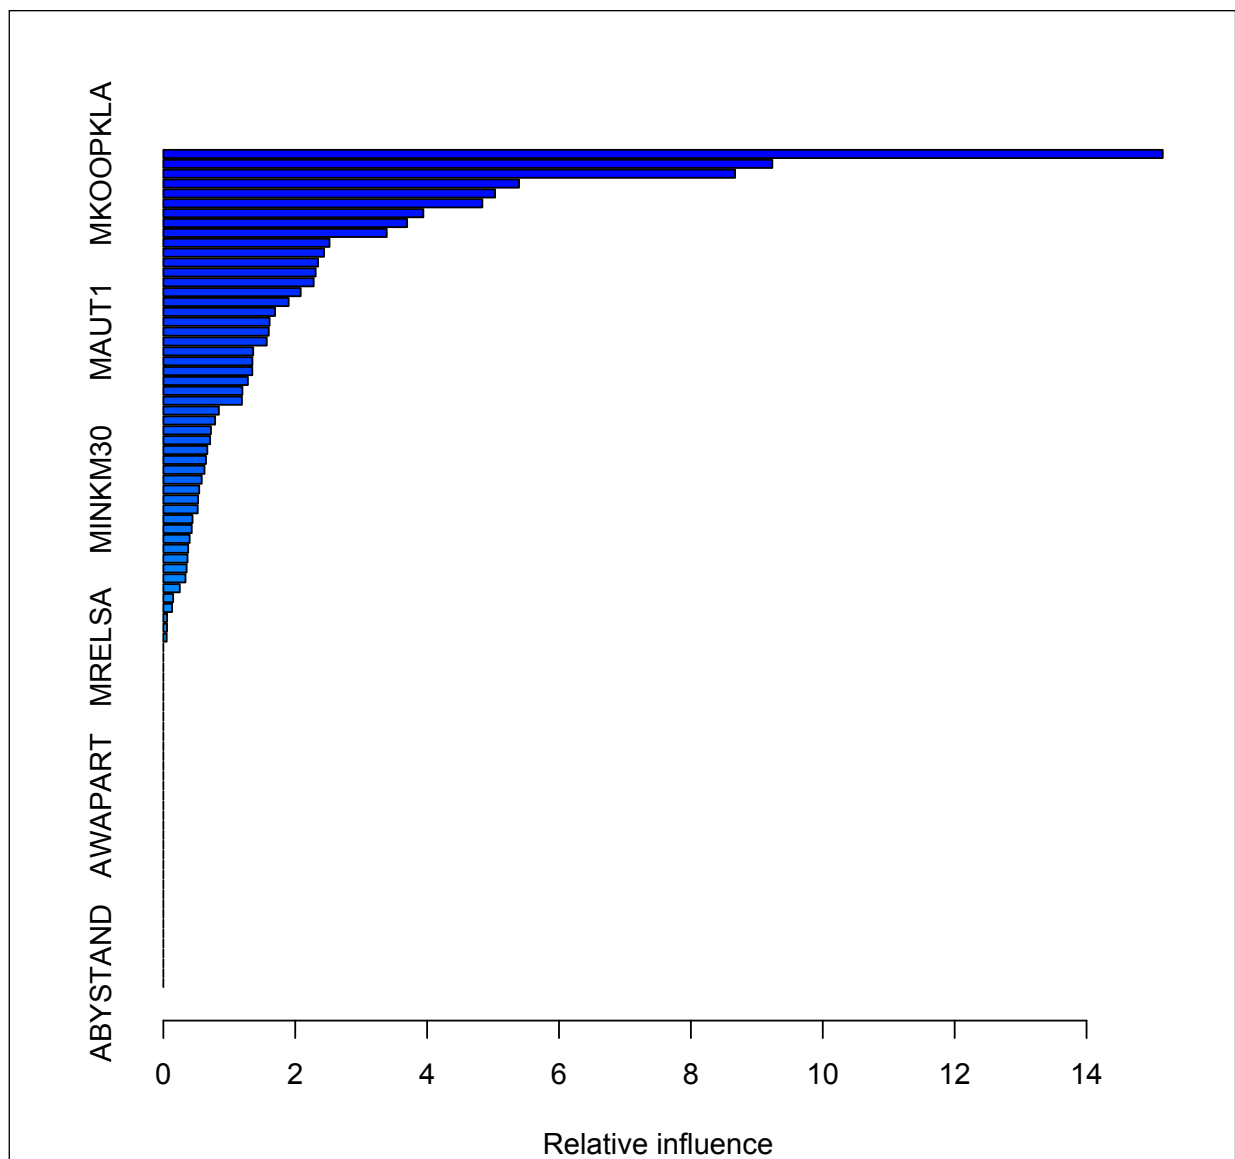
Answer:

```
> library(gbm)
> set.seed(342)
> boost.caravan = gbm(Purchase ~ ., data = Caravan.train, n.trees = 1000,
shrinkage = 0.01,
+   distribution = "bernoulli")
Warning messages:
1: In gbm.fit(x, y, offset = offset, distribution = distribution, w = w, :
  variable 50: PVRAAUT has no variation.
2: In gbm.fit(x, y, offset = offset, distribution = distribution, w = w, :
  variable 71: AVRAAUT has no variation.
> summary(boost.caravan)
```

| | var | rel.inf |
|----------|----------|-------------|
| PPERSAUT | PPERSAUT | 15.15534009 |
| MKOOPKLA | MKOOPKLA | 9.23499526 |
| MOPLHOOG | MOPLHOOG | 8.67017024 |
| MBERMIDD | MBERMIDD | 5.39403655 |
| MGODGE | MGODGE | 5.03047673 |
| PBRAND | PBRAND | 4.83740038 |
| MINK3045 | MINK3045 | 3.94305387 |
| ABRAND | ABRAND | 3.69692919 |
| MOSTYPE | MOSTYPE | 3.38768960 |
| PWAPART | PWAPART | 2.51970169 |
| MGODPR | MGODPR | 2.43689096 |
| MSKC | MSKC | 2.34594774 |
| MAUT2 | MAUT2 | 2.30973409 |
| MFEKIND | MFEKIND | 2.27959503 |
| MBERARBG | MBERARBG | 2.08245286 |
| MSKA | MSKA | 1.90020973 |
| PBYSTAND | PBYSTAND | 1.69481877 |
| MGODOV | MGODOV | 1.61147668 |
| MAUT1 | MAUT1 | 1.59879109 |
| MBERHOOG | MBERHOOG | 1.56791308 |
| MINK7512 | MINK7512 | 1.36255296 |
| MSKB1 | MSKB1 | 1.35071475 |

| | | |
|----------|----------|------------|
| MINKGEM | MINKGEM | 1.34913011 |
| MRELGE | MRELGE | 1.28204167 |
| MAUT0 | MAUT0 | 1.19929798 |
| MHHUUR | MHHUUR | 1.19158719 |
| MFGEKIND | MFGEKIND | 0.84203310 |
| MRELOV | MRELOV | 0.78554535 |
| MZPART | MZPART | 0.72191139 |
| MINK4575 | MINK4575 | 0.70935967 |
| MSKB2 | MSKB2 | 0.66694112 |
| APERSAUT | APERSAUT | 0.64644681 |
| MGODRK | MGODRK | 0.62380797 |
| MSKD | MSKD | 0.58168337 |
| MINKM30 | MINKM30 | 0.54392696 |
| PMOTSCO | PMOTSCO | 0.52708603 |
| MOPLMIDD | MOPLMIDD | 0.52091706 |
| MGEMOMV | MGEMOMV | 0.44231264 |
| MZFONDS | MZFONDS | 0.43037800 |
| PLEVEN | PLEVEN | 0.39901552 |
| MHKOOP | MHKOOP | 0.37672230 |
| MBERARBO | MBERARBO | 0.36653424 |
| MBERBOER | MBERBOER | 0.35290257 |
| MINK123M | MINK123M | 0.33559225 |
| MGEMLEEF | MGEMLEEF | 0.24937634 |
| MFALLEEN | MFALLEEN | 0.14898856 |
| MOSHOOFD | MOSHOOFD | 0.13265308 |
| MOPLLAAG | MOPLLAAG | 0.05654615 |
| MBERZELF | MBERZELF | 0.05589282 |
| MAANTHUI | MAANTHUI | 0.05047841 |
| MRELSA | MRELSA | 0.00000000 |
| PWABEDR | PWABEDR | 0.00000000 |
| PWALAND | PWALAND | 0.00000000 |
| PBESAUT | PBESAUT | 0.00000000 |
| PVRAAUT | PVRAAUT | 0.00000000 |
| PAANHANG | PAANHANG | 0.00000000 |
| PTRACTOR | PTRACTOR | 0.00000000 |
| PWERKT | PWERKT | 0.00000000 |
| PBROM | PBROM | 0.00000000 |
| PPERSONG | PPERSONG | 0.00000000 |
| PGEZONG | PGEZONG | 0.00000000 |
| PWAOREG | PWAOREG | 0.00000000 |
| PZEILPL | PZEILPL | 0.00000000 |
| PPLEZIER | PPLEZIER | 0.00000000 |
| PFIETS | PFIETS | 0.00000000 |
| PINBOED | PINBOED | 0.00000000 |
| AWAPART | AWAPART | 0.00000000 |
| AWABEDR | AWABEDR | 0.00000000 |

| | | |
|-----------|-----------|------------|
| AWALAND | AWALAND | 0.00000000 |
| ABESAUT | ABESAUT | 0.00000000 |
| AMOTSCO | AMOTSCO | 0.00000000 |
| AVRAAUT | AVRAAUT | 0.00000000 |
| AAANHANG | AAANHANG | 0.00000000 |
| ATTRACTOR | ATTRACTOR | 0.00000000 |
| AWERKT | AWERKT | 0.00000000 |
| ABROM | ABROM | 0.00000000 |
| ALEVEN | ALEVEN | 0.00000000 |
| APERSONG | APERSONG | 0.00000000 |
| AGEZONG | AGEZONG | 0.00000000 |
| AWAOREG | AWAOREG | 0.00000000 |
| AZEILPL | AZEILPL | 0.00000000 |
| APLEZIER | APLEZIER | 0.00000000 |
| AFIETS | AFIETS | 0.00000000 |
| AINBOED | AINBOED | 0.00000000 |
| ABYSTAND | ABYSTAND | 0.00000000 |



PPERSAUT, MKOOPKLA and MOPLHOOG are three most important variables in that order

- (c)

Answer:

```
> boost.pred = predict(boost.caravan, Caravan.test, n.trees = 1000,
type = "response")
> boost.pred = ifelse(boost.pred > 0.2, 1, 0)
> table(Caravan.test$Purchase, boost.pred)
    boost.pred
    0      1
0 4396 137
1  255  34
> 34/(137 + 34)
```

```
[1] 0.1988304
```

19.9% people predicted to make purchase will actually do.

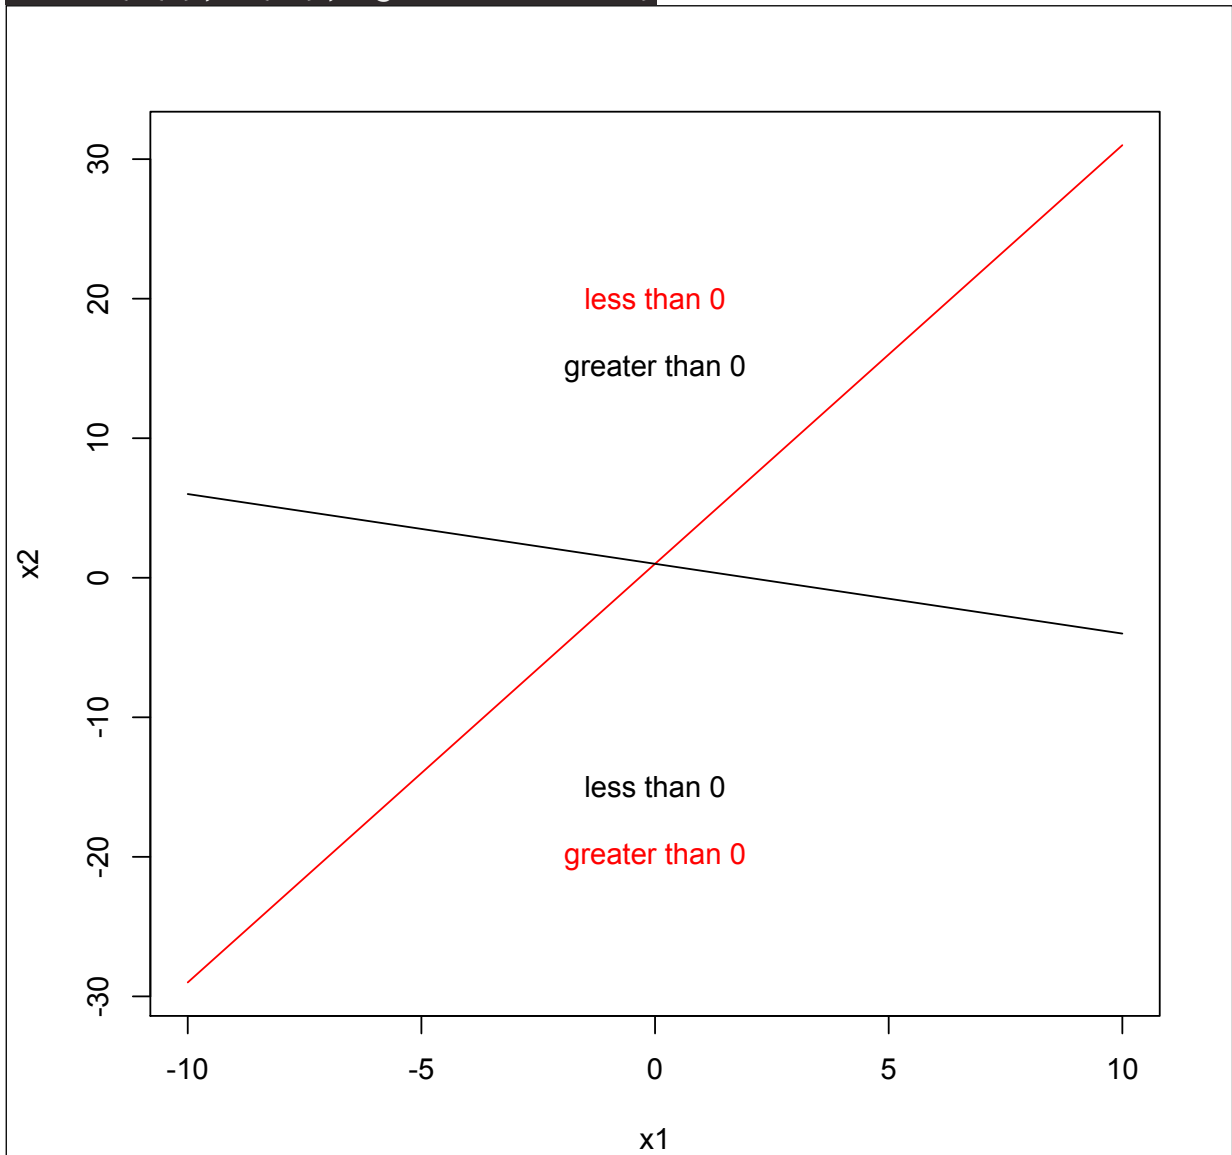
```
> lm.caravan = glm(Purchase ~ ., data = Caravan.train, family = binomial)
Warning message:
glm.fit: fitted probabilities numerically 0 or 1 occurred
> lm.prob = predict(lm.caravan, Caravan.test, type = "response")
Warning message:
In predict.lm(object, newdata, se.fit, scale = 1, type = ifelse(type == :
  prediction from a rank-deficient fit may be misleading
> lm.pred = ifelse(lm.prob > 0.2, 1, 0)
> table(Caravan.test$Purchase, lm.pred)
  lm.pred
      0      1
0 4183  350
1  231   58
> 58/(350 + 58)
[1] 0.1421569
```

Logistic regression about 14% people predicted to make purchase will do.

Problem 3 (p.368, ex1)

Answer:

```
> x1 = -10:10
> x2 = 1 + 3 * x1
> plot(x1, x2, type = "l", col = "red")
> text(c(0), c(-20), "greater than 0", col = "red")
> text(c(0), c(20), "less than 0", col = "red")
> lines(x1, 1 - x1/2)
> text(c(0), c(-15), "less than 0")
> text(c(0), c(15), "greater than 0")
```



Problem 4 (p.371, ex8)

(a)

Answer:

```
> library(ISLR)
> set.seed(9004)
> train = sample(dim(OJ)[1], 800)
> OJ.train = OJ[train, ]
> OJ.test = OJ[-train, ]
```

(b)

```
> library(e1071)
> svm.linear = svm(Purchase ~ ., kernel = "linear", data = OJ.train, cost = 0.01)
> summary(svm.linear)
```

```
Call:
svm(formula = Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.01)
```

```
Parameters:
  SVM-Type:  C-classification
  SVM-Kernel: linear
    cost:    0.01
   gamma:    0.05555556
```

```
Number of Support Vectors: 432
```

```
( 217 215 )
```

```
Number of Classes: 2
```

```
Levels:
CH MM
```

Answer:

SVC generates 432 support vectors out of 800 training points. 217 belong to level CH and 215 belong to level MM.

(c)

```
> train.pred = predict(svm.linear, OJ.train)
> table(OJ.train$Purchase, train.pred)
      train.pred
      CH  MM
CH 439  53
MM  82 226
> (82 + 53)/(439 + 53 + 82 + 226)
```

```
[1] 0.16875
> test.pred = predict(svm.linear, OJ.test)
> table(OJ.test$Purchase, test.pred)
  test.pred
    CH  MM
CH 142  19
MM  29  80
> (19 + 29)/(142 + 19 + 29 + 80)
[1] 0.1777778
```

Answer: The Training error rate is 16.9% and test error rate is 17.8%.

(d)

```
> set.seed(1554)
> tune.out = tune(svm, Purchase ~ ., data = OJ.train, kernel = "linear",
+ ranges = list(cost = 10^seq(-2,
+ 1, by = 0.25)))
summary(tune.out)
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost
0.3162278
```

- best performance: 0.16875

```
- Detailed performance results:
      cost error dispersion
1  0.01000000 0.16875 0.03691676
2  0.01778279 0.16875 0.03397814
3  0.03162278 0.17125 0.03230175
4  0.05623413 0.17250 0.03162278
5  0.10000000 0.17000 0.03291403
6  0.17782794 0.17125 0.03335936
7  0.31622777 0.16875 0.03498512
8  0.56234133 0.17000 0.03129164
9  1.00000000 0.16875 0.03397814
10 1.77827941 0.16875 0.03240906
11 3.16227766 0.16875 0.03294039
12 5.62341325 0.17125 0.03120831
13 10.00000000 0.17125 0.03283481
```

Answer:

We can see that the optimal cost is 0.3162278.

(e)

```
> svm.linear = svm(Purchase ~ ., kernel = "linear", data = OJ.train,
cost = tune.out$best.parameters$cost)
> train.pred = predict(svm.linear, OJ.train)
> table(OJ.train$Purchase, train.pred)
      train.pred
      CH  MM
CH 435  57
MM  71 237
> (57 + 71)/(435 + 57 + 71 + 237)
[1] 0.16
> test.pred = predict(svm.linear, OJ.test)
> table(OJ.test$Purchase, test.pred)
      test.pred
      CH  MM
CH 141  20
MM  29  80
> (29 + 20)/(141 + 20 + 29 + 80)
[1] 0.1814815
```

Answer:

Using the above cost, the training error decreases to 16% while test error increases to 18.1%.

(f)

```
> set.seed(410)
> svm.radial = svm(Purchase ~ ., data = OJ.train, kernel = "radial")
> summary(svm.radial)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "radial")
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1
gamma: 0.05555556
```

Number of Support Vectors: 367

```
( 184 183 )
```

```
Number of Classes: 2
```

```
Levels:
CH MM
```

```
> train.pred = predict(svm.radial, OJ.train)
> table(OJ.train$Purchase, train.pred)
  train.pred
    CH  MM
CH 452  40
MM  78 230
> (40 + 78)/(452 + 40 + 78 + 230)
[1] 0.1475
> test.pred = predict(svm.radial, OJ.test)
> table(OJ.test$Purchase, test.pred)
  test.pred
    CH  MM
CH 146  15
MM  27  82
> (27 + 15)/(146 + 15 + 27 + 82)
[1] 0.1555556
```

Answer:

We see that the radial basis kernel with default gamma creates 367 support vectors, out of which, 184 belong to level CH and remaining 183 belong to level MM. The classifier has a training error of 14.7% and a test error of 15.6% which is a slight improvement over linear kernel. To find optimal gamma:

```
> set.seed(755)
> tune.out = tune(svm, Purchase ~ ., data = OJ.train, kernel = "radial",
+   ranges = list(cost = 10^seq(-2,
+   1, by = 0.25)))
summary(tune.out)
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:
cost
0.5623413

- best performance: 0.165

- Detailed performance results:

| | cost | error | dispersion |
|----|-------------|---------|------------|
| 1 | 0.01000000 | 0.38500 | 0.06258328 |
| 2 | 0.01778279 | 0.38500 | 0.06258328 |
| 3 | 0.03162278 | 0.37625 | 0.06908379 |
| 4 | 0.05623413 | 0.21000 | 0.03855011 |
| 5 | 0.10000000 | 0.18625 | 0.03143004 |
| 6 | 0.17782794 | 0.18375 | 0.03230175 |
| 7 | 0.31622777 | 0.17125 | 0.03438447 |
| 8 | 0.56234133 | 0.16500 | 0.03763863 |
| 9 | 1.00000000 | 0.17500 | 0.03584302 |
| 10 | 1.77827941 | 0.17375 | 0.04059026 |
| 11 | 3.16227766 | 0.17625 | 0.03747684 |
| 12 | 5.62341325 | 0.17625 | 0.03839216 |
| 13 | 10.00000000 | 0.17375 | 0.03458584 |

```

> svm.radial = svm(Purchase ~ ., data = OJ.train, kernel = "radial", cost =
tune.out$best.parameters$cost)
> train.pred = predict(svm.radial, OJ.train)
> table(OJ.train$Purchase, train.pred)
      train.pred
      CH  MM
CH 452  40
MM  77 231
> (77 + 40)/(452 + 40 + 77 + 231)
[1] 0.14625
> test.pred = predict(svm.radial, OJ.test)
> table(OJ.test$Purchase, test.pred)
      test.pred
      CH  MM
CH 146  15
MM  28  81
> (28 + 15)/(146 + 15 + 28 + 81)
[1] 0.1592593

```

Tuning slightly decreases training error to 14.6% and increases test error to 16% which is still better than linear kernel.

(g)

```
> set.seed(8112)
> svm.poly = svm(Purchase ~ ., data = OJ.train, kernel = "poly", degree = 2)
> summary(svm.poly)
```

Call:

```
svm(formula = Purchase ~ ., data = OJ.train, kernel = "poly", degree = 2)
```

Parameters:

```
SVM-Type: C-classification
SVM-Kernel: polynomial
cost: 1
degree: 2
gamma: 0.05555556
coef.0: 0
```

Number of Support Vectors: 452

```
( 232 220 )
```

Number of Classes: 2

Levels:

```
CH MM
```

```
> train.pred = predict(svm.poly, OJ.train)
> table(OJ.train$Purchase, train.pred)
      train.pred
      CH  MM
CH 460  32
MM 105 203
> (12 + 37)/(149 + 12 + 37 + 72)
[1] 0.1814815
```

According to the summary, polynomial kernel produces 452 support vectors, out of which, 232 are level CH and 220 are level MM. This kernel produces a train error of 17.1% and a test error of 18.1% which are slightly higher than the errors produced by radial kernel but lower than the errors produced by linear kernel.

```
> set.seed(322)
> tune.out = tune(svm, Purchase ~ ., data = OJ.train, kernel = "poly",
degree = 2,
```

```
+ ranges = list(cost = 10^seq(-2, 1, by = 0.25)))
summary(tune.out)
> summary(tune.out)
```

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:

```
cost
5.623413
```

- best performance: 0.18375

- Detailed performance results:

| | cost | error | dispersion |
|----|-------------|---------|------------|
| 1 | 0.01000000 | 0.38500 | 0.05426274 |
| 2 | 0.01778279 | 0.36750 | 0.05075814 |
| 3 | 0.03162278 | 0.35750 | 0.05177408 |
| 4 | 0.05623413 | 0.34250 | 0.04937104 |
| 5 | 0.10000000 | 0.31500 | 0.05230785 |
| 6 | 0.17782794 | 0.24875 | 0.03928617 |
| 7 | 0.31622777 | 0.20875 | 0.05684103 |
| 8 | 0.56234133 | 0.20875 | 0.05653477 |
| 9 | 1.00000000 | 0.20000 | 0.06095308 |
| 10 | 1.77827941 | 0.19375 | 0.04497299 |
| 11 | 3.16227766 | 0.18625 | 0.04185375 |
| 12 | 5.62341325 | 0.18375 | 0.03335936 |
| 13 | 10.00000000 | 0.18375 | 0.04041881 |

```
> svm.poly = svm(Purchase ~ ., data = OJ.train, kernel = "poly", degree = 2,
cost = tune.out$best.parameters$cost)
```

```
> train.pred = predict(svm.poly, OJ.train)
```

```
> table(OJ.train$Purchase, train.pred)
```

```
      train.pred
      CH  MM
CH 455  37
MM  84 224
```

```
> (37 + 84)/(455 + 37 + 84 + 224)
```

```
[1] 0.15125
```

```
> test.pred = predict(svm.poly, OJ.test)
```

```
> table(OJ.test$Purchase, test.pred)
```

```
      test.pred
      CH  MM
CH 148  13
MM  34  75
```

```
> (13 + 34)/(148 + 13 + 34 + 75)
[1] 0.1740741
```

We can see that tuning reduces the training error to 15.12% and test error to 17.4% which is worse than radial kernel but better than linear kernel.

(h)

Radial basis kernel is the best since it produces minimum misclassification error on training and test data set.