

Lecture 1:
Course logistics,
Supervised vs. Unsupervised learning,
Bias-Variance tradeoff

STATS 202: Data mining and analysis

Rajan Patel

Syllabus

- ▶ **Videos:** Every lecture will be recorded by SCPD.
- ▶ **Email policy:** Please use the stats202 google group for most questions. Homeworks and all SCPD Exams should be e-mailed to stats202@gmail.com.
- ▶ **Class website:** www.stats202.com.

Prediction challenges

The MNIST dataset is a library of handwritten digits.



In a prediction challenge, you are given a **training** set of images of handwritten digits, which are labeled from 0 to 9.

You are also given a **test** set of handwritten digits, which are not identified.

Your job is to assign a digit to each image in the test set.

The Netflix prize

Netflix popularized prediction challenges by organizing an open, blind contest to improve its recommendation system.


The prize was \$1 million.



The Netflix prize

Netflix popularized prediction challenges by organizing an open, blind contest to improve its recommendation system.

The prize was \$1 million.




Some rankings were hidden in the training data

The Netflix prize

Netflix popularized prediction challenges by organizing an open, blind contest to improve its recommendation system.


The prize was \$1 million.



The challenge was to predict those rankings

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:



Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:



Variables or factors

Quantitative, eg. weight, height, number of children, ...

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:



Variables or factors

Qualitative, eg. college major, profession, gender, ...

Supervised vs. unsupervised learning

In **unsupervised learning** we start with a data matrix:


Our goal is to:

- ▶ Find meaningful relationships between the variables or units. **Correlation analysis.**
- ▶ Find meaningful groupings of the data. **Clustering.**

Unsupervised learning is also known in Statistics as **exploratory data analysis**.


Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:




Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:



Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:



Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

If X is the vector of inputs for a particular sample. The output variable is modeled by:

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Supervised vs. unsupervised learning

In **supervised learning**, there are *input* variables, and *output* variables:

If X is the vector of inputs for a particular sample. The output variable is modeled by:

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Our goal is to learn the function f , using a set of **training** samples.

Supervised vs. unsupervised learning

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Motivations:

- **Prediction:** Useful when the input variable is readily available, but the output variable is not.

Example: Predict stock prices next month using data from last year.

Supervised vs. unsupervised learning

$$Y = f(X) + \underbrace{\varepsilon}_{\text{Random error}}$$

Motivations:

- ▶ **Prediction:** Useful when the input variable is readily available, but the output variable is not.
- ▶ **Inference:** A model for f can help us understand the structure of the data — which variables influence the output, and which don't? What is the relationship between each variable and the output, e.g. linear, non-linear?

Example: What is the influence of genetic variations on the incidence of heart disease.

Parametric and nonparametric methods:

There are two kinds of supervised learning methods:

- ▶ **Parametric methods:** We assume that f takes a specific form. For example, a linear form:

$$f(X) = X_1\beta_1 + \cdots + X_p\beta_p$$

with parameters β_1, \dots, β_p . Using the training data, we try to *fit* the parameters.

- ▶ **Non-parametric methods:** We don't make any assumptions on the form of f , but we restrict how “wiggly” or “rough” the function can be.

Parametric vs. nonparametric prediction

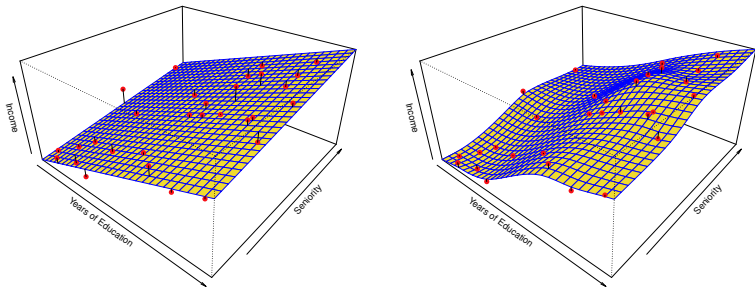


Figure: *

Figures 2.4 and 2.5

Parametric vs. nonparametric prediction




Figure: *

Figures 2.4 and 2.5

Parametric methods have a limit of fit quality. Non-parametric methods keep improving as we add more data to fit.

Parametric vs. nonparametric prediction

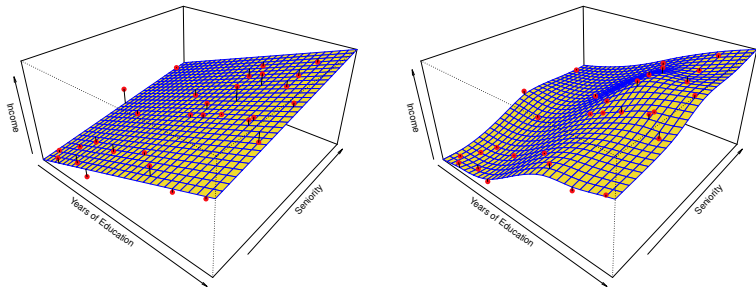


Figure: *

Figures 2.4 and 2.5

Parametric methods have a limit of fit quality. Non-parametric methods keep improving as we add more data to fit.

Parametric methods are often simpler to interpret.

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)$

Predicted function: \hat{f} .

Our goal in supervised learning is to minimize the prediction error.

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)$

Predicted function: \hat{f} .

Our goal in supervised learning is to minimize the **prediction error**.
For regression models, this is typically the *Mean Squared Error*:

$$MSE(\hat{f}) = E(y_0 - \hat{f}(x_0))^2.$$

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)$

Predicted function: \hat{f} .

Our goal in supervised learning is to minimize the **prediction error**.
For regression models, this is typically the *Mean Squared Error*:

$$MSE(\hat{f}) = E(y_0 - \hat{f}(x_0))^2.$$

Unfortunately, this quantity cannot be computed, because we don't know the true joint distribution of (X, Y) .

Prediction error

Training data: $(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)$

Predicted function: \hat{f} .

Our goal in supervised learning is to minimize the **prediction error**.
For regression models, this is typically the *Mean Squared Error*:

$$MSE(\hat{f}) = E(y_0 - \hat{f}(x_0))^2.$$

Unfortunately, this quantity cannot be computed, because we don't know the true joint distribution of (X, Y) . We can compute a sample average using the **training data**; this is known as the training MSE:

$$MSE_{\text{training}}(\hat{f}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2.$$

Prediction error

The main challenge of statistical learning is that *a low training MSE does not imply a low MSE.*

Prediction error

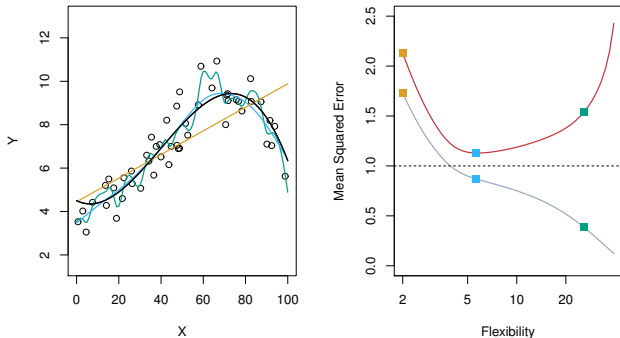
The main challenge of statistical learning is that *a low training MSE does not imply a low MSE.*

If we have test data $\{(x'_i, y'_i); i = 1, \dots, m\}$ which were not used to fit the model, a better measure of quality for \hat{f} is the test MSE:

$$MSE_{\text{test}}(\hat{f}) = \frac{1}{m} \sum_{i=1}^m (y'_i - \hat{f}(x'_i))^2.$$

Figure: *

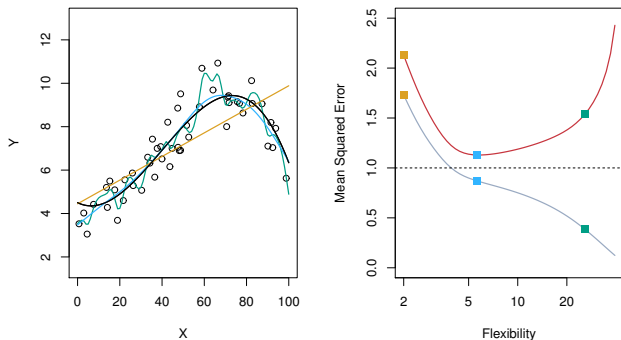
Figure 2.9.



The circles are simulated data from the black curve.

Figure: *

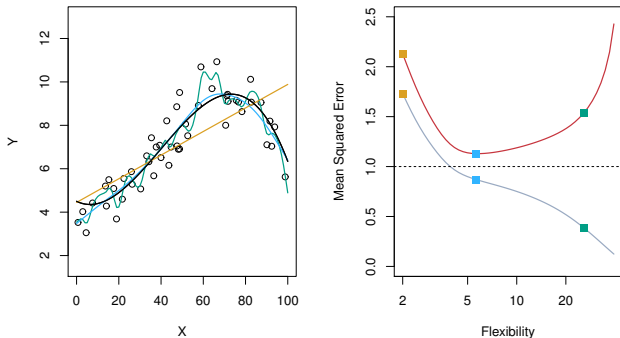
Figure 2.9.



The circles are simulated data from the black curve. In this artificial example, we *know* what f is.

Figure: *

Figure 2.9.

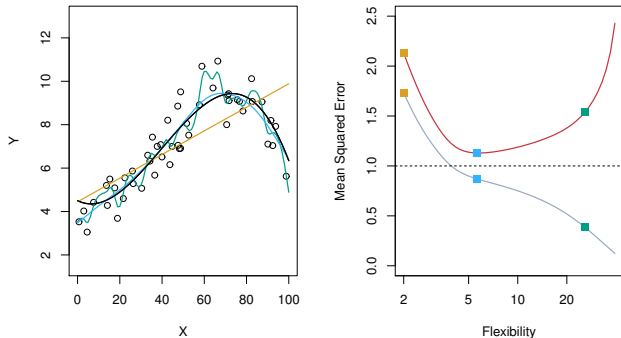


Three estimates \hat{f} are shown:

1. Linear regression.
2. Splines (very smooth).
3. Splines (quite rough).

Figure: *

Figure 2.9.



Red line: Test MSE.

Gray line: Training MSE.

Figure: *




Figure 2.10

The function f is now almost linear.

Figure: *




Figure 2.11

When the noise ε has small variance, the third method does well.

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and \hat{f} be estimated from n training samples $(x_1, y_1) \dots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs (y_1, \dots, y_n) . Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and \hat{f} be estimated from n training samples $(x_1, y_1) \dots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs (y_1, \dots, y_n) . Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

Irreducible error

The bias variance decomposition

Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and \hat{f} be estimated from n training samples $(x_1, y_1) \dots (x_n, y_n)$.

Let E denote the expectation over y_0 and the training outputs (y_1, \dots, y_n) . Then, the Mean Squared Error at x_0 can be decomposed:

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$

The variance of the estimate of Y : $E[\hat{f}(x_0) - E(\hat{f}(x_0))]^2$

This measures how much the estimate of \hat{f} at x_0 changes when we sample new training data.

The bias variance decomposition


Let x_0 be a fixed test point, $y_0 = f(x_0) + \varepsilon_0$, and \hat{f} be estimated from n training samples $(x_1, y_1) \dots (x_n, y_n)$.



Let E denote the expectation over y_0 and the training outputs (y_1, \dots, y_n) . Then, the Mean Squared Error at x_0 can be decomposed:




$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon_0).$$





The squared bias of the estimate of Y : $[E(\hat{f}(x_0)) - f(x_0)]^2$





This measures the deviation of the average prediction $\hat{f}(x_0)$ from the truth $f(x_0)$.













































Implications of bias variance decomposition

$$MSE(x_0) = E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\varepsilon).$$


- ▶ The MSE is always positive.
- ▶ Each element on the right hand side is always positive.
- ▶ Therefore, typically when we decrease the bias beyond some point, we increase the variance, and vice-versa.

More flexibility \iff Higher variance \iff Lower bias.

Squiggly f , high noise



Linear f , high noise



Squiggly f , low noise




Figure: *

Figure 2.12

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz, } \dots\}$.

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz, } \dots\}$.

The model:

$$Y = f(X) + \varepsilon$$

becomes insufficient, as X is not necessarily real-valued.

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz, } \dots\}$.

The model:

$$\cancel{Y = f(X) + \varepsilon}$$

becomes insufficient, as X is not necessarily real-valued.

Classification problems

In a classification setting, the output takes values in a discrete set.

For example, if we are predicting the brand of a car based on a number of variables, the function f takes values in the set $\{\text{Ford, Toyota, Mercedes-Benz, } \dots\}$.

We will use slightly different notation:

$P(X, Y)$: joint distribution of (X, Y) ,
 $P(Y | X)$: conditional distribution of X given Y ,
 \hat{y}_i : prediction for x_i .

Loss function for classification

There are many ways to measure the error of a classification prediction. One of the most common is the 0-1 loss:

$$E(\mathbf{1}(y_0 \neq \hat{y}_0))$$

Like the MSE, this quantity can be estimated from training and test data by taking a sample average:

$$\frac{1}{n} \sum_{i=1}^n \mathbf{1}(y_i \neq \hat{y}_i)$$